

Lecture A4. Simulation 1

Sim, Min Kyu, Ph.D., mksim@seoultech.ac.kr



서울과학기술대학교 데이터사이언스학과

1 I. Motivation - estimation of π

2 II. Simulation approach

3 III. Discussion

4 IV. Confidence interval

I. Motivation - estimation of π

What is π ?

- π is defined as

$$\pi = \frac{\text{a circle's circumference}}{\text{a circle's diameter}}$$

- To list a few reasons why π is such an important quantity:
 - In Architecture
 - In Construction
 - In Art
 - In Military operation
 - so many...

How to estimate?

- In your elementary school
- From your high school, you learned that

$$\begin{aligned} \text{(quarter size of a unit circle)} &= \int_0^1 \sqrt{1-x^2} dx \\ &= \pi/4 \end{aligned}$$

- In ancient days, people used $\pi/4 = \int_0^1 \sqrt{1-x^2} dx$ to estimate π .

II. Simulation approach

Design

- Step 1.
 - Let $X \sim U(-1, 1)$ and $Y \sim U(-1, 1)$.
 - Generate two vectors length of N , i.e. $\mathbf{x} = (x_1, x_2, \dots, x_N)$ and $\mathbf{y} = (y_1, y_2, \dots, y_N)$, where x_i is a sample of X and y_i is a sample of Y for all i .
- Step 2.
 - Let $t_i := \sqrt{x_i^2 + y_i^2}$ for all i , that is,

$$t_1 := \sqrt{x_1^2 + y_1^2}$$

$$t_2 := \sqrt{x_2^2 + y_2^2}$$

$$\dots \quad \dots$$

$$t_N := \sqrt{x_N^2 + y_N^2}$$

- Step 3.

$$\hat{\pi} = 4 \times \frac{\text{number of } \{t_i \leq 1\}}{N} = 4 \times \frac{\sum_{i=1}^N I_{\{t_i \leq 1\}}}{N}$$

Remark 1

The function $I_{\{.\}}$ is called an *indicator function* that returns 1 if the statement is true and 0 if false.

- $\frac{\sum_{i=1}^N I_{\{t_i \leq 1\}}}{N}$ counts the number of t_i that is less than or equal to 1, among all i .

A statistical software, R

- You should comfortably interchange between R and python.
- Resources 1 - 'R for Python user'
 - <https://medium.com/@nawazahmad20/r-for-python-programmers-part-1-ca4eab668b8c>
 - <http://ramnathv.github.io/pycon2014-r/>
- Resources 2 - datacamp.com
 - datacamp.com allows access for >100 courses for R and python.
 - I suggest you at least do 'Introduction to R' and 'Intermediate R'.
 - You can subscribe for free using your @seoultech.ac.kr email with the following link.
 - https://www.datacamp.com/groups/shared_links/b3b5fc6f798aaf54ada0c03cee875c0099c34e300f5be6b8375e4850646b0b59
 - The above link expires every March and September, but I always renew it.
 - You can visit my github and find the link <https://github.com/aceMKSIm/teaching/>
 - You can always request me for an invitation link if yours is expired.
- Resources 3 - lecture material for data visualization
 - From the following repository, study L01-L03 for installation and basic usage.
 - <https://github.com/aceMKSIm/teaching/tree/master/Data%20Visualization/Lecture%20Notes>

Implementation - basic

- Implementation with 1000 repetitions.

```
set.seed(1234) # fix the random seed
MC_N <- 10^3
x <- runif(MC_N)*2-1 # runif() generates  $U(0,1)$ 
y <- runif(MC_N)*2-1 # this code generates  $U(-1,1)$ 
t <- sqrt(x^2+y^2)
head(cbind(x,y,t)) # always display and check!
```

```
##           x           y           t
## [1,] -0.7725932  0.6752678  1.0261028
## [2,]  0.2445988 -0.0250675  0.2458800
## [3,]  0.2185495 -0.7793260  0.8093904
## [4,]  0.2467589 -0.2972740  0.3863441
## [5,]  0.7218308  0.5221261  0.8908733
## [6,]  0.2806212 -0.2206703  0.3569925
```

```
pi_hat <- 4*sum(t<=1)/MC_N
pi_hat
```

```
## [1] 3.188
```

- 1 `set.seed()` fixes randomization, which is often convenient to get consistent outcome.
- 2 `runif(MC_N)` generates a vector of length `MC_N`, where each element follows $U(0, 1)$.
- 3 `runif(MC_N)*2` follows $U(0, 2)$, and `runif(MC_N)*2-1` follows $U(-1, 1)$.
- 4 `t<=1` returns the 0-1 vector of length same as `t`, where an element is 1 if corresponding element in `t` is less than or equal to 1, and 0 otherwise.
- 5 `cbind()` combines (column) vectors into a matrix.
- 6 `head()` displays the first six observations.

Vectorized programming

- From the previous slide

```
beg_time <- Sys.time()
set.seed(1234)
MC_N <- 10^6
x <- runif(MC_N)*2-1
y <- runif(MC_N)*2-1
t <- sqrt(x^2+y^2)
pi_hat <- 4*sum(t<=1)/MC_N
end_time <- Sys.time()
print(end_time-beg_time)

## Time difference of 0.07878995 secs
```

- What a first-timer would write.

```
beg_time <- Sys.time()
set.seed(1234)
MC_N <- 10^6
count <- 0
for (MC_i in 1:MC_N) {
  x_i <- runif(1)*2-1
  y_i <- runif(1)*2-1
  t_i <- sqrt(x_i^2+y_i^2)
  if (t_i <= 1) count <- count + 1
}
pi_hat <- 4*count/MC_N
end_time <- Sys.time()
print(end_time-beg_time)
```

Time difference of 2.858543 secs

- The style of the code on the left is called *vectorized programming*.
- It is elegant, economic, and efficient.
- You must be able to *write as the left side* and *communicate as the right side* (to non-expert).

Exercise 1

Write the two programs in the previous page with python and compare the computation time.

Implementation - varying number of trials

- Approach with a custom function

```
pi_simulator <- function(MC_N) {  
  set.seed(1234)  
  x <- runif(MC_N)*2-1  
  y <- runif(MC_N)*2-1  
  t <- sqrt(x^2+y^2)  
  pi_hat <- 4*sum(t<=1)/MC_N  
  return(pi_hat)  
}
```

```
pi_simulator(100)
```

```
## [1] 3.04
```

```
pi_simulator(1000)
```

```
## [1] 3.188
```

```
pi_simulator(10000)
```

```
## [1] 3.1876
```

```
pi_simulator(100000)
```

```
## [1] 3.13432
```

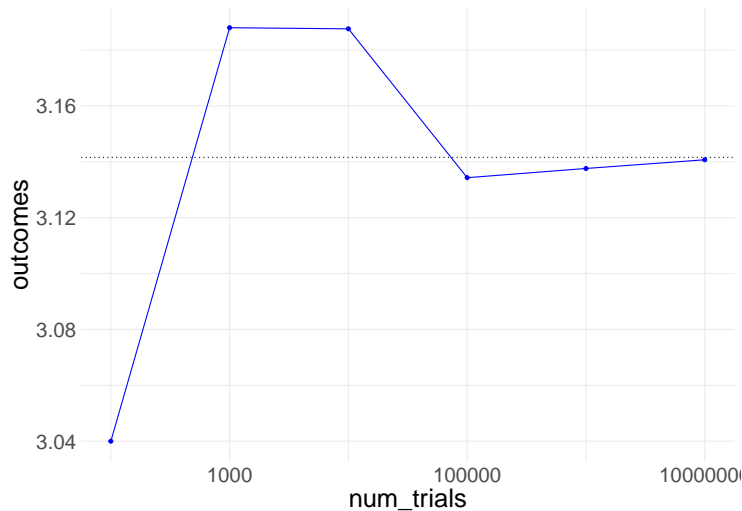
- How many repetition is necessary to get closer?

```
num_trials <- 10^(2:7)  
outcomes <- sapply(num_trials, pi_simulator)  
results <- cbind(num_trials, outcomes)  
results
```

```
##      num_trials outcomes  
## [1,]         100 3.040000  
## [2,]        1000 3.188000  
## [3,]       10000 3.187600  
## [4,]      100000 3.134320  
## [5,]     1000000 3.137616  
## [6,]    10000000 3.140733
```

- `sapply(num_trials, pi_simulator)`
applies the function `pi_simulator()`
to *each element* of `num_trials`.

- How many repetition is necessary to get closer?



- (optional) The previous figure was plotted by the following code.

```
results <- data.frame(results)
library(tidyverse)
ggplot(results, aes(x=num_trials, y=outcomes)) +
  geom_point(color = "blue") + geom_path(color = "blue") +
  geom_abline(slope = 0, intercept = 3.14159, linetype = "dotted") +
  scale_x_log10() +
  theme_minimal() + theme(text = element_text(size=25))
```


III. Discussion

Computation time

- In this implementation, number of trials were increased from 10^2 to 10^7 . No wonder that this increases the computational time.
- Following modified function displays the elapsed time.


```
pi_simulator2 <- function(MC_N) { # name change
  beg_time <- Sys.time() # newly added
  set.seed(1234)
  x <- runif(MC_N)*2-1
  y <- runif(MC_N)*2-1
  t <- sqrt(x^2+y^2)
  pi_hat <- 4*sum(t<=1)/MC_N
  end_time <- Sys.time() # newly added
  print(MC_N)
  print(end_time-beg_time) # newly added
  return(pi_hat)
}
```

```
sapply(10^(2:6), pi_simulator2)
```

```
## [1] 100
## Time difference of 0 secs
## [1] 1000
## Time difference of 0 secs
## [1] 10000
## Time difference of 0.0009980202 secs
## [1] 100000
## Time difference of 0.005984068 secs
## [1] 1000000
## Time difference of 0.07583094 secs
## [1] 3.040000 3.188000 3.187600 3.134320 3.137616
```

Confidence on the result

- In estimation of π example, we were in the luxurious situation because we already knew the correct value of π , 3.14159.
- In reality, this situation is rare. Rather, you shouldn't be in need of doing simulation after all if you already know the exact value.
- In reality, following figure is what you would normally face. Notice that the correct value indicating line is gone.



- One legitimate way to have some confidence is to observe the plot and say ‘It seems converging with good degree’. Then, present a number that seems to be within the tolerance.
- Are there any way to present a confidence interval, just as good statistical estimations should present?

IV. Confidence interval

Motivation

- Building a confidence interval from experiment generally involves repetitive experiments. But in the simulation approach, we already do have the repetitive simulation experiments? Isn't this enough?
- Not really so. In order to build a confidence interval, we should treat *one entire simulation experiment* as *one observation*. For example, we treat the result from a MC_N=1000 simulation experiment as a single observation on the true value. Then repeat this simulation experiment, say, n times, to build a confidence interval.
- Let's set MC_N=1000 for a simulation experiment and do this for $n=100$ times.

Repetitive simulation experiments

- Let's set $MC_N=1,000$ for a simulation experiment and do this for $n=100$ times.
- For each experiment, record the result to collect $n=100$ samples.

```
pi_simulator3 <- function(MC_N) { # name change
  # set.seed(1234) # seed must not be fixed
  x <- runif(MC_N)*2-1
  y <- runif(MC_N)*2-1
  t <- sqrt(x^2+y^2)
  pi_hat <- 4*sum(t<=1)/MC_N
  return(pi_hat)
}
```

```
n <- 100 # number of experiments to repeat
MC_N <- 1000 # number of simulation repetition in a single experiment
set.seed(1234)
samples <- rep(0, n) # create an empty zero vector
for (i in 1:n) { # do this for n times
  samples[i] <- pi_simulator3(MC_N)
}
head(samples)
```

```
## [1] 3.188 3.144 3.060 3.240 3.148 3.172
```

- From LN.A4.p13,

$$\mathbb{P}[\bar{X} - t_{0.975, n-1} \cdot s/\sqrt{n} \leq \mu \leq \bar{X} + t_{0.975, n-1} \cdot s/\sqrt{n}] = 0.95$$

- Obtain the numbers as follows:

```
X_bar <- mean(samples)
s <- sqrt(sum((X_bar-samples)^2)/(n-1))
t <- qt(p=0.975, df = n-1)
```

```
X_bar
## [1] 3.137
s
## [1] 0.05186579
t
## [1] 1.984217
```

- Thus,

$$\mathbb{P}[3.137 - 1.984 \cdot 0.0519/\sqrt{100} \leq \mu \leq 3.137 + 1.984 \cdot 0.0519/\sqrt{100}] = 0.95$$

\Downarrow

$$\mathbb{P}[3.127 \leq \mu \leq 3.147] = 0.95$$

- Note that the length of interval was 0.020 (=3.147-3.127)
- Obviously, increasing MC_N and/or increasing n should narrow the confidence interval.

Exercise 2

Do the above experiment with MC_N increased by the factor of ten, and present the confidence interval. (Use `set.seed(1234)`)

```
n <- 100 # number of exp. to rep.
MC_N <- 10000 # number of sim. rep. in a single exp.
set.seed(1234)
samples <- rep(0, n)
for (i in 1:n) {
  samples[i] <- pi_simulator3(MC_N)
}
X_bar <- mean(samples)
s <- sqrt(sum((X_bar-samples)^2)/(n-1))
t <- qt(p=0.975, df = n-1)
lb <- X_bar-t*s/sqrt(n) # Lower bound
ub <- X_bar+t*s/sqrt(n) # upper bound
```

```
lb
## [1] 3.137164
ub
## [1] 3.143788
ub-lb
## [1] 0.006624215
```

Exercise 3

Do the Exercise above with n increased by the factor of ten, and present the confidence interval. (Use `set.seed(1234)`)

```
n <- 1000 # number of exp. to rep.
MC_N <- 10000 # number of sim. rep. in a single exp.
set.seed(1234)
samples <- rep(0, n)
for (i in 1:n) {
  samples[i] <- pi_simulator3(MC_N)
}
X_bar <- mean(samples)
s <- sqrt(sum((X_bar-samples)^2)/(n-1))
t <- qt(p=0.975, df = n-1)
lb <- X_bar-t*s/sqrt(n) # Lower bound
ub <- X_bar+t*s/sqrt(n) # upper bound
```

```
lb
## [1] 3.139777
ub
## [1] 3.141834
ub-lb
## [1] 0.002057237
```

Computation cost and The accuracy

MC_N	n	length of CI
1,000	100	0.020
1,0000	100	0.0066
1,0000	1000	0.00205

- Increasing MC_N or n gives the same effect.
 - When MC_N was increased by the factor of 10, the length of CI was decreased by the factor of $\sqrt{10}$.
 - When n was increased by the factor of 10, the length of CI was decreased by the factor of $\sqrt{10}$.
- Repetitive simulation experiments is beneficial if...
 - when you need confidence interval.
 - when you face memory issue that prevents increasing MC_N any more.

Exercise 4

Present a previous page's table by writing a neat python code block.

"If I only had an hour to chop down a tree, I would spend the first 45 minutes sharpening my axe.
- A. Lincoln"